

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

In [2]:

```
df = pd.read_csv('Telco-Customer-Churn.csv')
```

In [3]:

```
df.head()
```

Out[3]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	N
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Y
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Y
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Y
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	N

5 rows x 21 columns



In [4]:

```
# Confirm that there are no NaN cells by displaying NaN values per feature column.
df.isna().sum()
```

Out[4]:

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

In [5]:

```
#df.dropna()
```

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   customerID            7032 non-null   object  
1   gender                 7032 non-null   object  
2   SeniorCitizen          7032 non-null   int64   
3   Partner                7032 non-null   object  
4   Dependents             7032 non-null   object  
5   tenure                 7032 non-null   int64   
6   PhoneService           7032 non-null   object  
7   MultipleLines          7032 non-null   object  
8   InternetService        7032 non-null   object  
9   OnlineSecurity         7032 non-null   object  
10  OnlineBackup           7032 non-null   object  
11  DeviceProtection       7032 non-null   object  
12  TechSupport            7032 non-null   object  
13  StreamingTV            7032 non-null   object  
14  StreamingMovies        7032 non-null   object  
15  Contract               7032 non-null   object  
16  PaperlessBilling       7032 non-null   object  
17  PaymentMethod          7032 non-null   object  
18  MonthlyCharges         7032 non-null   float64  
19  TotalCharges           7032 non-null   float64  
20  Churn                  7032 non-null   object  
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

In [7]:

```
df.describe()
```

Out[7]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7032.000000	7032.000000	7032.000000	7032.000000
mean	0.162400	32.421786	64.798208	2283.300441
std	0.368844	24.545260	30.085974	2266.771362
min	0.000000	1.000000	18.250000	18.800000
25%	0.000000	9.000000	35.587500	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.862500	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

In [8]:

```
df['Churn'].value_counts()
```

Out[8]:


```
No      5163
Yes     1869
Name: Churn, dtype: int64
```

In [9]:

```
plt.figure(figsize=(15,8), dpi=200)
```

```
sns.countplot(data = df, x = 'Churn')
```

```
plt.savefig('fig1.png')
```




In [10]:

```
plt.figure(figsize=(15,8), dpi = 200)
```

```
sns.violinplot(data = df, x = 'Churn' , y = 'TotalCharges')
```

```
plt.savefig('fig2.png')
```





In [11]:

```
plt.figure(figsize=(15,8), dpi = 200)
```

```
sns.boxplot(data = df, y = 'TotalCharges', x = 'Contract', hue = 'Churn')
```

```
plt.savefig('fig3.png')
```





In [12]:

```
df.columns
```

Out[12]:

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

In [13]:

```
corr_df = pd.get_dummies(df[['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
                              'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
                              'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
                              'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
                              'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn']]).corr()
```

In [14]:

```
corr_df
```

Out[14]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	customerID_0002-ORFBO	customerID_0003-MKNFE	cus
SeniorCitizen	1.000000	0.015683	0.219874	0.102411	-0.005251	-0.005251	
tenure	0.015683	1.000000	0.246862	0.825880	-0.011381	-0.011381	
MonthlyCharges	0.219874	0.246862	1.000000	0.651065	0.000318	-0.001942	
TotalCharges	0.102411	0.825880	0.651065	1.000000	-0.008892	-0.009160	
customerID_0002-ORFBO	-0.005251	-0.011381	0.000318	-0.008892	1.000000	-0.000142	
customerID_0003-MKNFE	-0.005251	-0.011381	-0.001942	-0.009160	-0.000142	1.000000	
PaymentMethod_Credit card (automatic)	-0.024359	0.232800	0.030055	0.182663	-0.006265	-0.006265	
PaymentMethod_Electronic check	0.171322	0.210197	0.271117	-0.060436	-0.008490	-0.008490	
PaymentMethod_Mailed check	-0.152987	0.232181	-0.376568	-0.294708	0.021939	0.021939	

Churn_No	-0.150541	0.354049	-0.192858	0.199484	customerID_00025	customerID_00065	cus
SeniorCitizen	tenure	MonthlyCharges	TotalCharges	ORFBO	MKNFE		
Churn_Yes	0.150541	0.354049	0.192858	-0.199484	-0.007175	-0.007175	

7079 rows x 7079 columns

In [15]:

```
corr_df['Churn_Yes'].sort_values()
```

Out[15]:

```
Churn_No          -1.000000
tenure           -0.354049
Contract_Two year -0.301552
StreamingMovies_No internet service -0.227578
StreamingTV_No internet service -0.227578
...
InternetService_Fiber optic 0.307463
TechSupport_No          0.336877
OnlineSecurity_No       0.342235
Contract_Month-to-month 0.404565
Churn_Yes              1.000000
Name: Churn_Yes, Length: 7079, dtype: float64
```

In [16]:

```
corr_df['Churn_Yes'].sort_values().iloc[1:-1]
```

Out[16]:

```
tenure           -0.354049
Contract_Two year -0.301552
StreamingMovies_No internet service -0.227578
StreamingTV_No internet service -0.227578
InternetService_No -0.227578
...
PaymentMethod_Electronic check 0.301455
InternetService_Fiber optic 0.307463
TechSupport_No          0.336877
OnlineSecurity_No       0.342235
Contract_Month-to-month 0.404565
Name: Churn_Yes, Length: 7077, dtype: float64
```

Churn Analysis

Now we focus on segmentating customers based on their tenure, creating 'cohorts', allowing us to examine difference between customer cohort segments

In [17]:

```
df['tenure'].unique()
```

Out[17]:

```
array([ 1, 34,  2, 45,  8, 22, 10, 28, 62, 13, 16, 58, 49, 25, 69, 52, 71,
        21, 12, 30, 47, 72, 17, 27,  5, 46, 11, 70, 63, 43, 15, 60, 18, 66,
         9,  3, 31, 50, 64, 56,  7, 42, 35, 48, 29, 65, 38, 68, 32, 55, 37,
        36, 41,  6,  4, 33, 67, 23, 57, 61, 14, 20, 53, 40, 59, 24, 44, 19,
        54, 51, 26, 39])
```


Create a histogram displaying the distribution of 'tenure' column, which is the amount of months a customer was or has been on a customer.

In [18]:

```
plt.figure(figsize=(10,4),dpi=200)
```

```
sns.histplot(data=df,x='tenure',bins=60)
```

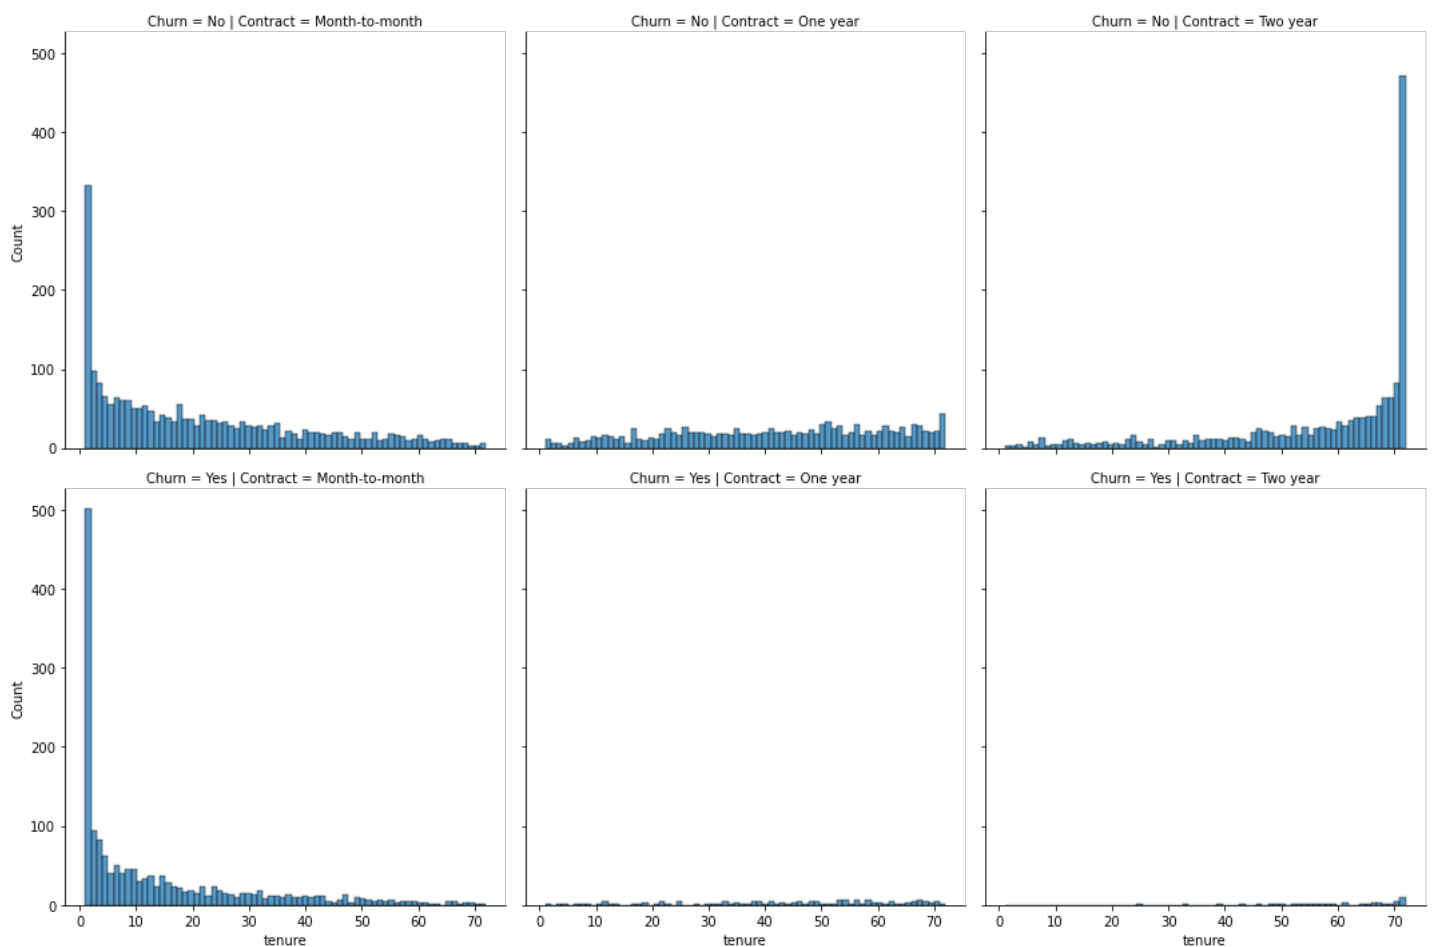
```
plt.savefig('fig4.png')
```



In [19]:

```
plt.figure(figsize=(10,3),dpi=200)
sns.displot(data=df,x='tenure',bins=70,col='Contract',row='Churn');
plt.savefig('fig5.png')
```

<Figure size 2000x600 with 0 Axes>




Display a scatter plot of Total Charges versus Monthly Charges, and color hue by Churn.

In [22]:

```
plt.figure(figsize=(10,3),dpi=200)
```

```
sns.scatterplot(data = df, x = 'MonthlyCharges', y = 'TotalCharges', hue = 'Churn',
                linewidth=0.5,alpha=0.5)
plt.savefig('fig6.png')
```



Creating Cohorts based on Tenure

Let's begin by treating each unique tenure length, 1 month, 2 month, 3 month...N months as its own cohort.

In [36]:

```
yes_churn = df.groupby(['Churn' , 'tenure']).count().transpose()['Yes']
no_churn = df.groupby(['Churn', 'tenure']).count().transpose()['No']
```

In [37]:

```
churn_rate = 100 * yes_churn / (no_churn+yes_churn)
```

In [39]:

```
churn_rate.transpose()
```

Out[39]:

	customerID	gender	SeniorCitizen	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSec
tenure									
1	61.990212	61.990212	61.990212	61.990212	61.990212	61.990212	61.990212	61.990212	61.990212
2	51.680672	51.680672	51.680672	51.680672	51.680672	51.680672	51.680672	51.680672	51.680672
3	47.000000	47.000000	47.000000	47.000000	47.000000	47.000000	47.000000	47.000000	47.000000
4	47.159091	47.159091	47.159091	47.159091	47.159091	47.159091	47.159091	47.159091	47.159091
5	48.120301	48.120301	48.120301	48.120301	48.120301	48.120301	48.120301	48.120301	48.120301
...
68	9.000000	9.000000	9.000000	9.000000	9.000000	9.000000	9.000000	9.000000	9.000000
69	8.421053	8.421053	8.421053	8.421053	8.421053	8.421053	8.421053	8.421053	8.421053
70	9.243697	9.243697	9.243697	9.243697	9.243697	9.243697	9.243697	9.243697	9.243697
71	3.529412	3.529412	3.529412	3.529412	3.529412	3.529412	3.529412	3.529412	3.529412
72	1.657459	1.657459	1.657459	1.657459	1.657459	1.657459	1.657459	1.657459	1.657459

72 rows x 19 columns

In [42]:

```
##churn_rate.transpose()['customerID'].plot()
```

```
In [43]:
df['tenure']

Out[43]:
0      1
1     34
2      2
3     45
4      2
..
7027   24
7028   72
7029   11
7030    4
7031   66
Name: tenure, Length: 7032, dtype: int64
```

Based on the tenure column values, create a new column called Tenure Cohort that creates 4 separate categories:

'0-12 Months'

'24-48 Months'

'12-24 Months'

'Over 48 Months'

```
In [49]:
def cohort(tenure):
    if tenure < 13:
        return '0-12 Months'
    elif tenure < 25:
        return '12-24 Months'
    elif tenure < 49:
        return '24-48 Months'
    else:
        return "Over 48 Months"
```

```
In [51]:
cohort(17)
```

Out[51]:
'12-24 Months'

```
In [53]:
df ['Tenure Cohort'] = df['tenure'].apply(cohort)
```

```
In [55]:
df.head(10)[['tenure', 'Tenure Cohort']]

Out[55]:
```


	tenure	Tenure Cohort
0	1	0-12 Months
1	34	24-48 Months
2	2	0-12 Months

3	tenure	Tenure Cohort
4	2	0-12 Months
5	8	0-12 Months
6	22	12-24 Months
7	10	0-12 Months
8	28	24-48 Months
9	62	Over 48 Months

In [64]:

```
## Now create a scatter plot between monthly charge and total charge
plt.figure(figsize=(14,6), dpi = 200)
sns.scatterplot(data = df, x = 'MonthlyCharges', y = 'TotalCharges', hue = 'Tenure Cohort',
                linewidth = 0.5, alpha = 0.5)

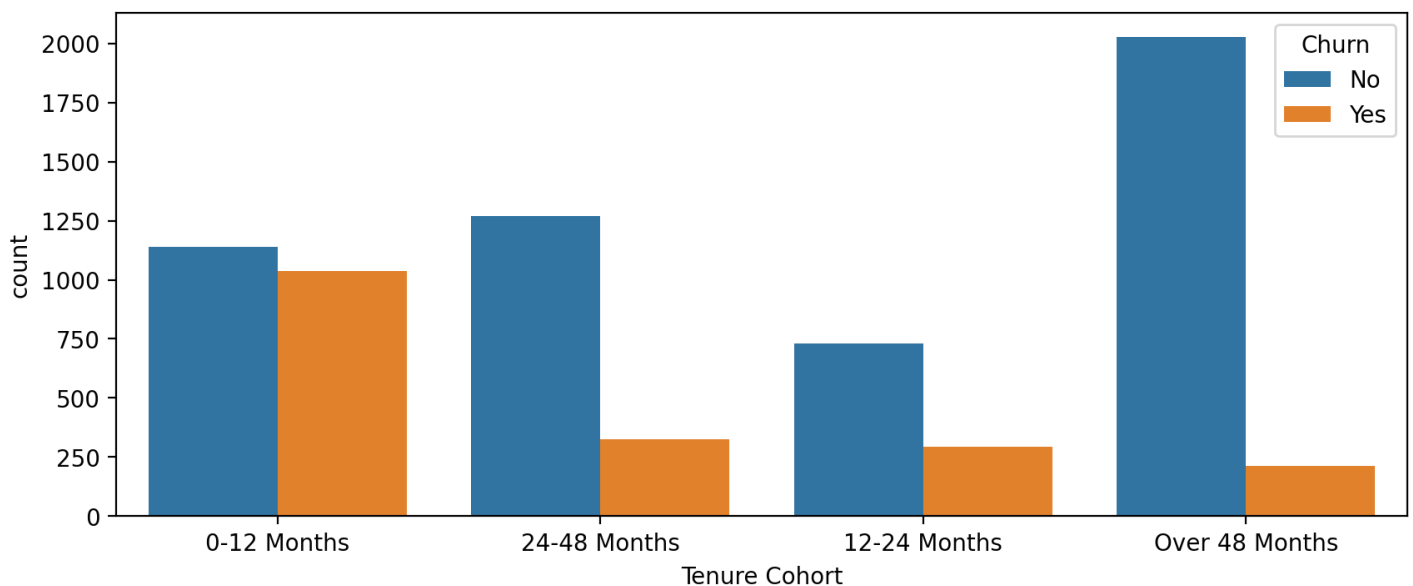
plt.savefig('fig7.png')
```



In [65]:

```
## Create a count plot the churn count per cohort
plt.figure(figsize=(10,4), dpi=200)
sns.countplot(data=df, x='Tenure Cohort', hue='Churn')

plt.savefig('fig8.png')
```



Now Lets discuss our Predictive Modeling

Let's explore 4 different tree based methods: A Single Decision Tree, Random Forest, AdaBoost, Gradient Boosting.

In [70]:

```
df.head(10)
```

Out[70]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No
5	9305-CDSKC	Female	0	No	No	8	Yes	Yes	Fiber optic	No
6	1452-KIOVK	Male	0	No	Yes	22	Yes	Yes	Fiber optic	No
7	6713-OKOMC	Female	0	No	No	10	No	No phone service	DSL	Yes
8	7892-POOKP	Female	0	Yes	No	28	Yes	Yes	Fiber optic	No
9	6388-TABGU	Male	0	No	Yes	62	Yes	No	DSL	Yes

10 rows x 22 columns



In [72]:

```
X = df.drop(['Churn' , 'customerID'],axis= 1)
```

In [73]:

```
X = pd.get_dummies(X,drop_first=True)
```

In [74]:

```
y = df['Churn']
```

In [75]:

```
from sklearn.model_selection import train_test_split
```

In [76]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 101)
```

In [77]:

```
## Decision Tree Performance
from sklearn.tree import DecisionTreeClassifier
```

In [78]:

```
dt = DecisionTreeClassifier(max_depth = 8)
```

In [79]:

```
dt.fit(X_train, y_train)
```

Out[79]:

```
DecisionTreeClassifier(max_depth=8)
```

In [80]:

```
pred = dt.predict(X_test)
```

In [82]:

```
from sklearn.metrics import plot_confusion_matrix, classification_report
```

In [84]:

```
print(classification_report(y_test, pred))
```


	precision	recall	f1-score	support
No	0.87	0.87	0.87	557
Yes	0.51	0.50	0.51	147
accuracy			0.80	704
macro avg	0.69	0.69	0.69	704
weighted avg	0.79	0.80	0.79	704

In [85]:

```
plot_confusion_matrix(dt, X_test, y_test)
```

Out[85]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe6ad830d00>
```




In [86]:

```
from sklearn.tree import plot_tree
```

```
In [88]:
```

```
plt.figure(figsize=(12,8),dpi=150)
plot_tree(dt,filled=True,feature_names=X.columns)

plt.savefig('fig9.png')
```



Random Forest

Now determine aRandom Forest model and create a classification report and confusion matrix from its predicted results on the test set.**

```
In [89]:
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
In [95]:
```

```
rf = RandomForestClassifier(n_estimators= 100, max_depth=8)
```

```
In [96]:
```

```
rf.fit(X_train,y_train)
```

```
Out[96]:
```

```
RandomForestClassifier(max_depth=8)
```

```
In [97]:
```

```
predts = rf.predict(X_test)
```

```
In [98]:
```

```
print(classification_report(y_test,predts))
```

	precision	recall	f1-score	support
No	0.87	0.92	0.89	557


Yes	0.61	0.47	0.53	147
accuracy			0.83	704
macro avg	0.74	0.69	0.71	704
weighted avg	0.81	0.83	0.82	704

In [112]:

```
plot_confusion_matrix(rf,X_test,y_test)
```

Out[112]:

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe6b1844790>



Boosted Trees

We are trying Ada Boost and gredient Boost to create a model and report back the classification report and plot a confusion matrix for its predicted results

In [114]:

```
from sklearn.ensemble import GradientBoostingClassifier,AdaBoostClassifier
```

In [115]:

```
ada_model = AdaBoostClassifier()
```

In [116]:

```
gb_model = GradientBoostingClassifier()
```

In [117]:

```
ada_model.fit(X_train,y_train)
```

Out[117]:

AdaBoostClassifier()

In [118]:

```
gb_model.fit(X_train,y_train)
```

Out[118]:

GradientBoostingClassifier()

In [119]:

```
ada_preds = ada_model.predict(X_test)
```

In [120]:

```
gb_preds = gb_model.predict(X_test)
```

In [121]:

```
print(classification_report(y_test,ada_preds))
```

	precision	recall	f1-score	support
No	0.88	0.90	0.89	557
Yes	0.60	0.54	0.57	147
accuracy			0.83	704
macro avg	0.74	0.72	0.73	704
weighted avg	0.82	0.83	0.83	704

In [122]:

```
print(classification_report(y_test,gb_preds))
```


	precision	recall	f1-score	support
No	0.87	0.90	0.89	557
Yes	0.57	0.50	0.53	147
accuracy			0.82	704
macro avg	0.72	0.70	0.71	704
weighted avg	0.81	0.82	0.81	704

In [123]:

```
plot_confusion_matrix(ada_model,X_test,y_test)
```

Out[123]:

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe6ad7e8400>





In [124]:

```
plot_confusion_matrix(gb_model,X_test,y_test)
```

Out[124]:

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe6ad7e80a0>





With base models, we got best performance from an AdaBoostClassifier, but note, we didn't do any gridsearching AND most models performed about the same on the data set.

Great Job!!

In []: